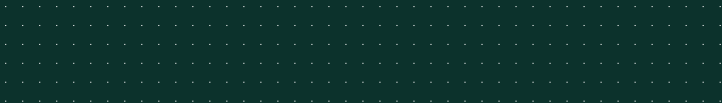


APRIL 25, OSTRAVA

warewulf making cluster installations fast and reliable

Christian Goll <cgoll@suse.com>



Introduction

warewulf is tool for managing beowulf clusters

Beowulf

- old british poem

Beowulf cluster

- became popular in the 90.
- use of the shelf hardware
 - 486 & linux
 - **not** Cray & unix
- warewulf is a typo of werewolf



Introduction

HPC landscape

Top five Supercomputers

1	Frontier	EPYC 64C	AMD MI250X	Slingshot-11
2	Aurora	Xeon 9470	Intel GPU Max	Slingshot-11
3	Eagle	Xeon 8480	NVIDIA H100	NVIDIA Infiniband
4	Fugaku	A64FX 48C 2.2GHz	-	Tofu interconnect D
5	LUMI	EPYC 64C 2GHz	AMD MI250X	Slingshot-11

- only Fugaku uses non standard CPU
- others are beowulf clusters with GPUs attached



Introduction

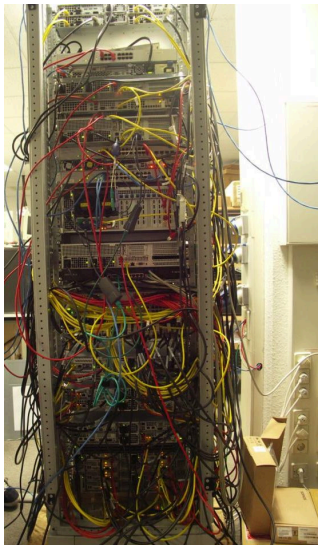
Beowulf cluster

base components

- management node
- compute nodes
- management network

optional components

- more compute nodes
- fast network interconnects
- central storage
- bmc/ipmi

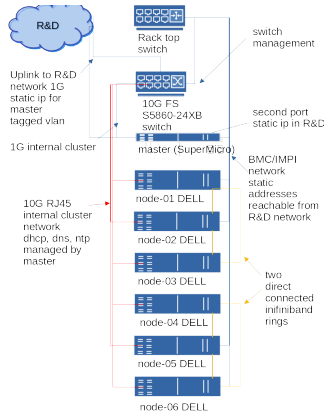


Introduction

Beowulf Cluster

differences to data centers

- compute nodes are cattle
- hierarchical organization
- compute are not updated after boot process
- application come from central storage
- applications are self compiled
- one application can run over several nodes



Warewulf description

software stack

warewulf components

warewulfd delivers

- kernel & modules
- node image
- node configurations

wwctl cmd line tool

- manages node database
- manages node image

external components

dhcp server

- ISC dhcpd server
- dnsmasq

tftp

- tftp from kernel.org
- dnsmasq

optional

- nfs
- manage /etc/hosts



Warewulf configuration

database /etc/warewulf/nodes.conf

- plain yaml file
- easy backup
- can be version controlled
- external tools support
 - vim, ansible

profiles

- stores identical values for collection of nodes
- values can be overridden on node basis

```
1 WW_INTERNAL: 45
2 nodeprofiles:
3   default:
4     comment: This profile is automatically
5     container name: leap
6     network devices:
7       default:
8         device: eth0
9 nodes:
10  n01:
11    profiles:
12    - default
13    network devices:
14      default:
15        hwaddr: 52:54:00:4e:cb:1d
16        ipaddr: 172.16.130.101
17  n02:
```



Warewulf configuration

command line database manipulation

add node

```
1 wwctl node add n01  
  -I 10.10.10.1
```

modify node

```
1 wwctl node set n01  
  --comment "Have fun"
```

list node

```
1 wwctl node list  
  n01 -a
```

	NODE	FIELD	PROFILE	VALUE
2	n01	Id	--	n01
3	n01	Comment	SUPERSEDED	Have fun
4	n01	ContainerName	default	leap
5	n01	Ipxe	--	(default)
6	n01	RuntimeOverlay	--	(generic)
7	n01	SystemOverlay	--	(wwinit)
8	n01	Root	--	(initramfs)
9	n01	Discoverable	--	false
10	n01	Init	--	(/sbin/init)
11	n01	Kernel.Args	--	(quiet crashkernel=)
12	n01	Profiles	--	default
13	n01	PrimaryNetDev	--	(default)
14	n01	NetDevs[default].Type	--	(ethernet)
15	n01	NetDevs[default].OnBoot	--	(true)
16	n01	NetDevs[default].Device	default	eth0
17	n01	NetDevs[default].Hwaddr	--	52:54:00:4e:cb:1d
	n01	NetDevs[default].Ipaddr	--	172.16.100.101



Warewulf configuration

templates& overlays

Configuration templates

- based on go templates
- `{{.foo}}` replaced with variable `foo`
- exported go function can be called

Configuration overlays

- rendered templates packed into overlay
- overlay put ontop of node image

Listing 1: /etc/issue.ww

```
1 Warewulf Node:      {{.Id}}
2 Container:          {{.Container}}
3 {{ if .Kernel.Version }}Kernel:
4 Kernelargs:          {{.Kernel.Args}}
5
6 Network:
7 {{- range $devname, $netdev := .NetDevs}}
8     {{$devname}}: {{$netdev.Device}}
9     {{$devname}}: {{$netdev.IpCIDR}}
10 {{if $netdev.Ipaddr6 }}      {{$devname}}:
11 {{if $netdev.Hwaddr }}      {{$devname}}:
12 {{end}}
```



Warewulf configuration

rendered templates

Configuration templates

- based on go templates
- `{{.foo}}` replaced with variable `foo`
- exported go function can be called

Configuration overlays

- rendered templates packed into overlay
- overlay put on top of node image

Listing 2: `/etc/issue.ww`

```
1 Warewulf Node:      {{.Id}}
2 Container:          {{.Container}}
3 {{ if .Kernel.Version }}Kernel:
4 Kernelargs:          {{.Kernel.Args}}
5
6 Network:
7 {{- range $devname, $netdev := .NetDevs}}
8     {{$devname}}: {{$netdev.Device}}
9     {{$devname}}: {{$netdev.IpCIDR}}
10 {{if $netdev.Ipaddr6 }}      {{$devname}}:
11 {{if $netdev.Hwaddr }}      {{$devname}}:
12 {{end}}
```



Warewulf configuration

overlays

warewulf defines two types of overlays

system overlay

- available on boot
- warewulf boot strap files
- static network configurations:
 - wicked
 - NetworkManager
 - EL scripts
- nfs mounts
- file system mounts

runtime overlay

- updated on regular base
- can be secured

user defined overlays

- users are encouraged to create own configuration templates
- can reside in system & runtime overlays



Warewulf configuration


security

assumptions

- private/cluster network is secure
- lateral movement isn't accounted
 - NFS mounts are common, when not **mandatory**

measurements

- node image & system overlay protected with BIOS UUID
- system overlays **must** be downloaded from privileged port



Warewulf configuration

node images

definition

- complete OS images
- called containers in warewulf
- must be imported from:
 - chroot directory
 - docker registry
 - local `dockerd`
- several different node images can be imported
- node images are vendor independent

registry.suse.com

- SUSE SLE 15SP5

registry.opensuse.org

- openSUSE Tumbleweed
- openSUSE Leap 15SP[3-5]

ghcr.io

- openSUSE Leap
- Rocky EL (8&9)
- Debian Bockworm



Warewulf configuration

node image examples

Import the SLE image

```
1 ww4-host:~> export WAREWULF_OCI_USERNAME=cgoll@suse.com
2 ww4-host:~> export WAREWULF_OCI_PASSWORD=INTERNAL-USE-ONLY-xxxxxx
3 ww4-host:~> wwctl container import docker://registry.suse.com/suse/hpc/
  warewulf4-x86_64/sle-hpc-node:latest sle-hpc
```

Import the Leap image

```
1 ww4-host:~> wwctl container import docker://registry.opensuse.org/science
  /warewulf/leap-15.5/containers/kernel:latest leap
```



Warewulf configuration

node image examples

Execute shell in images

```
1 wwctl container shell sle-hpc
2 WARN    : Couldn't mount /etc/SUSEConnect to /etc/SUSEConnect: no such
  file or directory
3 WARN    : Couldn't mount /etc/zypp/credentials.d/SCCcredentials to /etc/
  zypp/credentials.d/SCCcredentials: no such file or directory
4 [sle-hpc] Warewulf>
```

- SLE registration from outer node is mounted into image



Warewulf configuration

disk management

Needs following elements:

- disks
- partitions needs parent disk
- filesystem needs partent partition

implementation

- call `ignition` with `ignition-ww4-disk.service`
- **not** in `dracut`
- **before** `sysroot.mount`

single partition

```
1 wwctl node set n01 --diskname /dev/  
  vda --diskwipe --partname scratch  
    --partcreate --fsname scratch --  
    fsformat btrfs --fspath /scratch  
    --fswipe
```

add swap

```
1 wwctl node set n01 --diskname /dev/  
  vda --partname swap --partsize  
    =1024 --partnumber 1 --fsname  
    swap --fsformat swap --fspath  
    swap
```

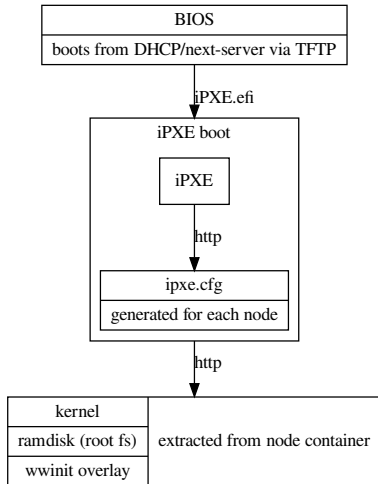


Warewulf boot

boot process

boot with iPXE

- distribution iPXE binaries are used
- `tftp` transfers are small
- `kernel` is extracted from container/node image on the fly
- root fs is the container/node image
- configuration overlay added on top
- **no** secure boot




Warewulf boot

boot process

boot with grub tftp

- grub & shim extracted from host
- secure boot only possible if host shim & grub can boot node kernel




Warewulf boot

boot process

boot with grub http

- grub & shim xtracted from node/container image
- secure boot with various distributions
- must be configured in BIOS



Warewulf

development

availability

- part of SLE since 15SP5
- SLE node/container available

upstream

github.com/warewulf/warewulf

Rocky Linux Foundation project

Stakeholders:

- SUSE
- CIQ
- Intel/openHPC



Thank your, for you attention!

